

Brothers in Arms

How to Make MySQL and PostgreSQL Work Together

Charly Batista

Senior Support Engineer

Percona Live 2018

2018-04-24



Agenda

- Introduction
- Replication
 - *PostgreSQL Replication*
 - *MySQL Replication*
- Challenges
- Available solutions
 - *SymmetricDS*
 - *Tungsten Replicator*
 - *Apache Kafka with Debezium*
- Demo

Introduction

Why replicate data between different databases?

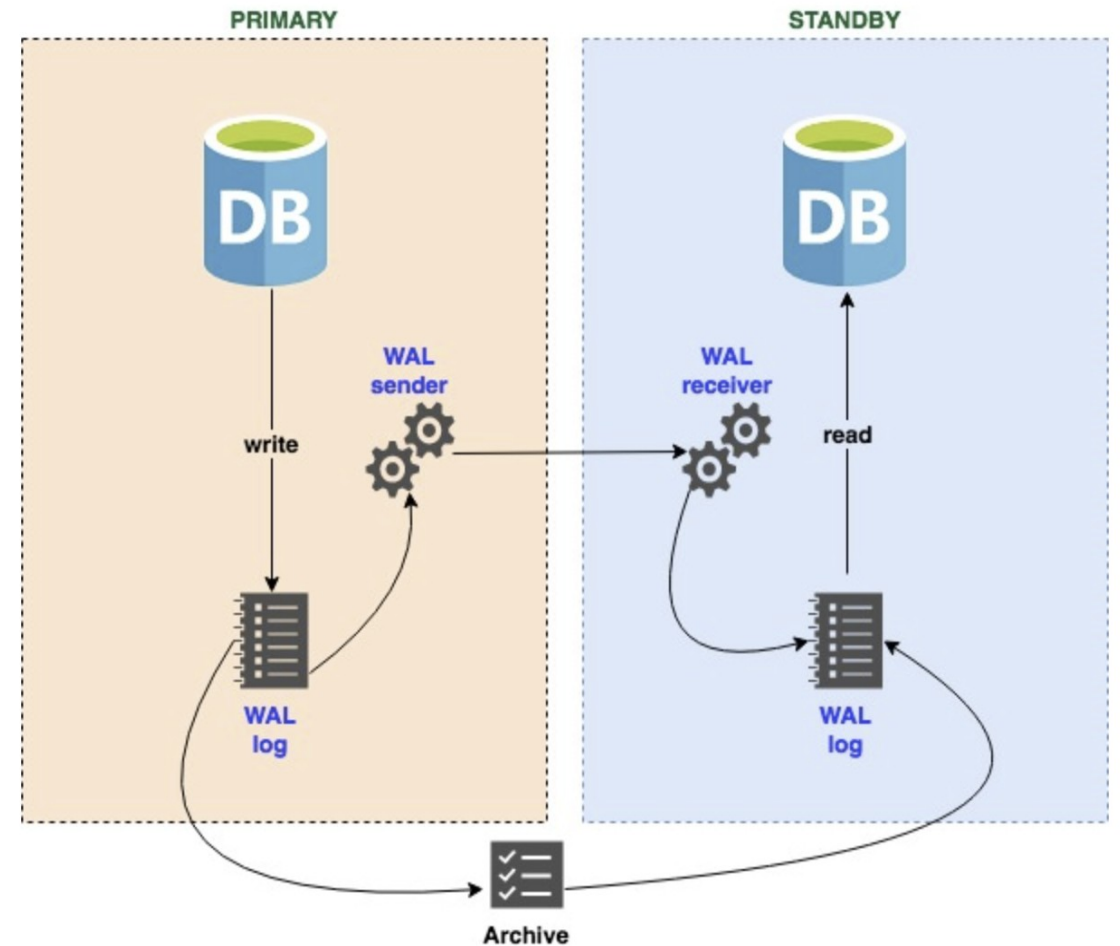
- Migrate data from one database to the another with minimal downtime;
- Feed legacy/in-house systems with web transactions;
- DW and Reporting;
- Use the correct technology assets (GEO, Sync Replication, json, etc);

Replication

Replication

PostgreSQL Streaming Replication

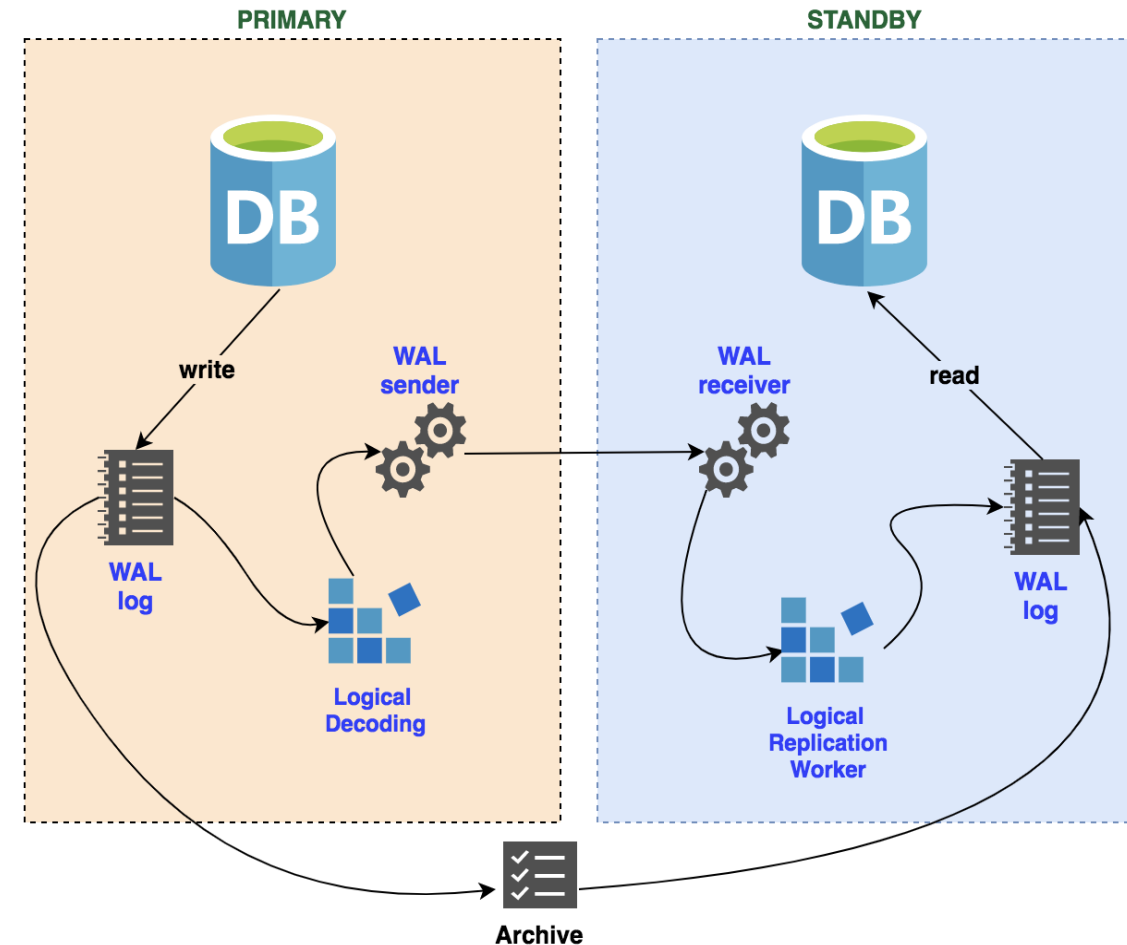
- Physical replication (blocks, byte-by-byte);
- Not possible with different major versions;
- Slaves readable, not writeable;
- Slaves have same structure as master;
- Complex timeline control;
- Re-sync slave often requires new BKP;



Replication...

PostgreSQL Logical Replication (PG 10+)

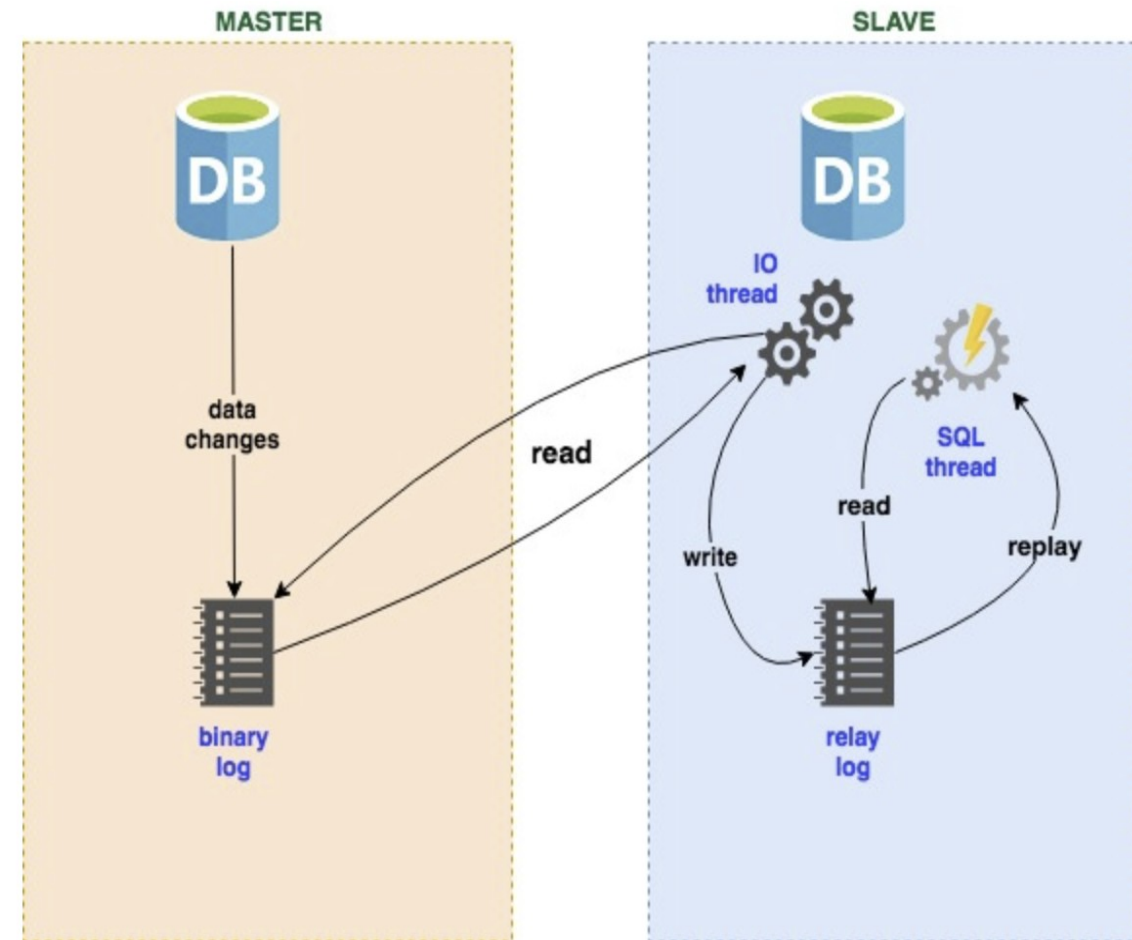
- Uses a publish and subscribe model;
- Slaves can now be writable;
- Now possible with different major versions;
- Transactional consistency is guaranteed for publications within a single subscription;
- Better granularity and control;
- Doesn't replicate DDL;



Replication...

MySQL

- Logical Replication;
- Slaves can be writable;
- Possible with different major versions*;
- Good granularity and control;
- Replicates schema and data;

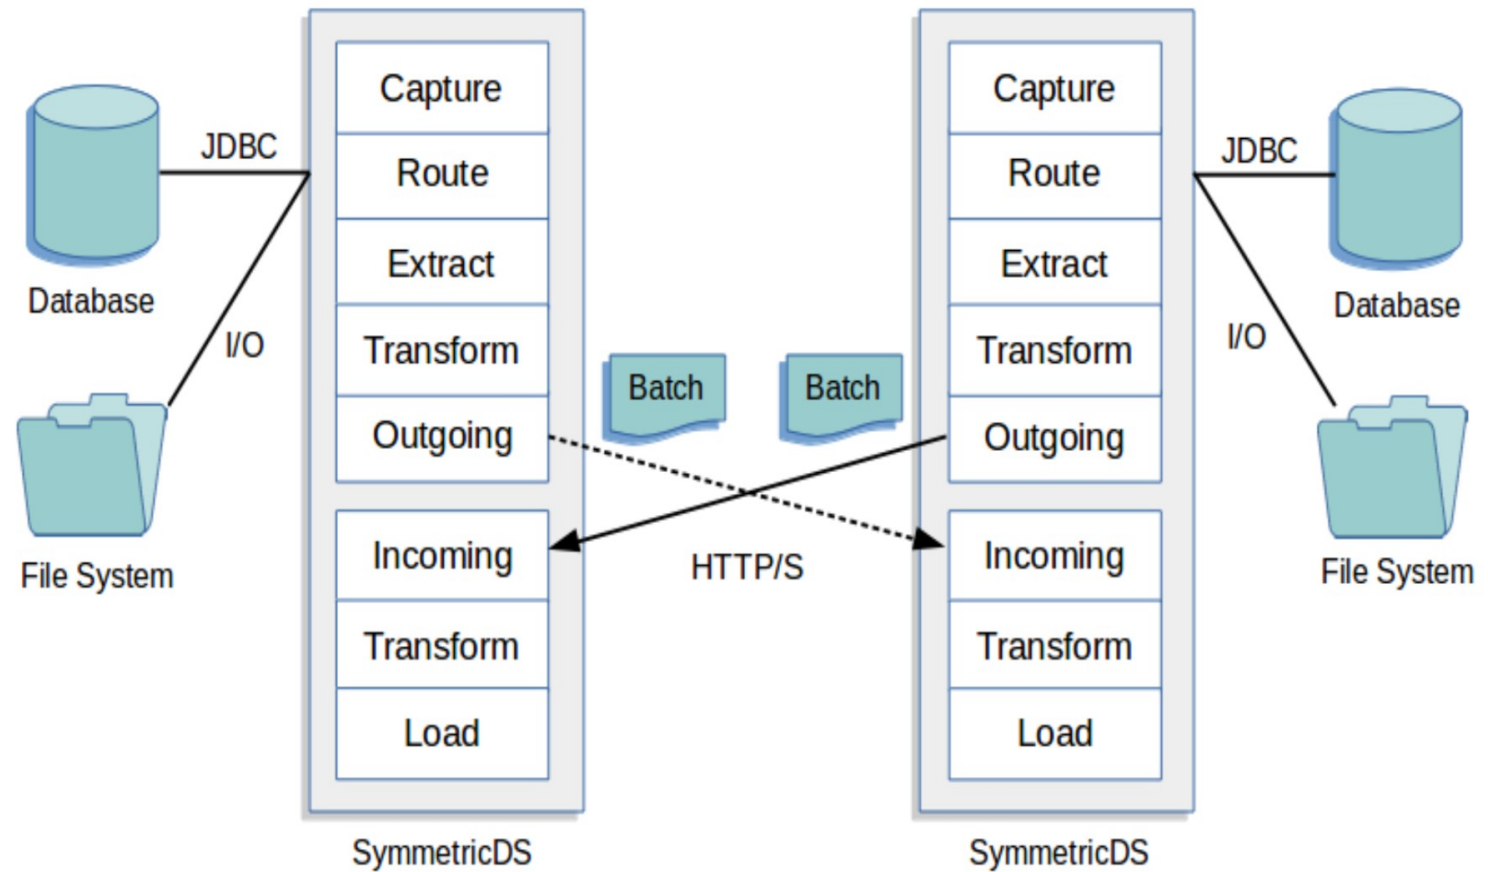


Available solutions

Available solutions

SymmetricDS

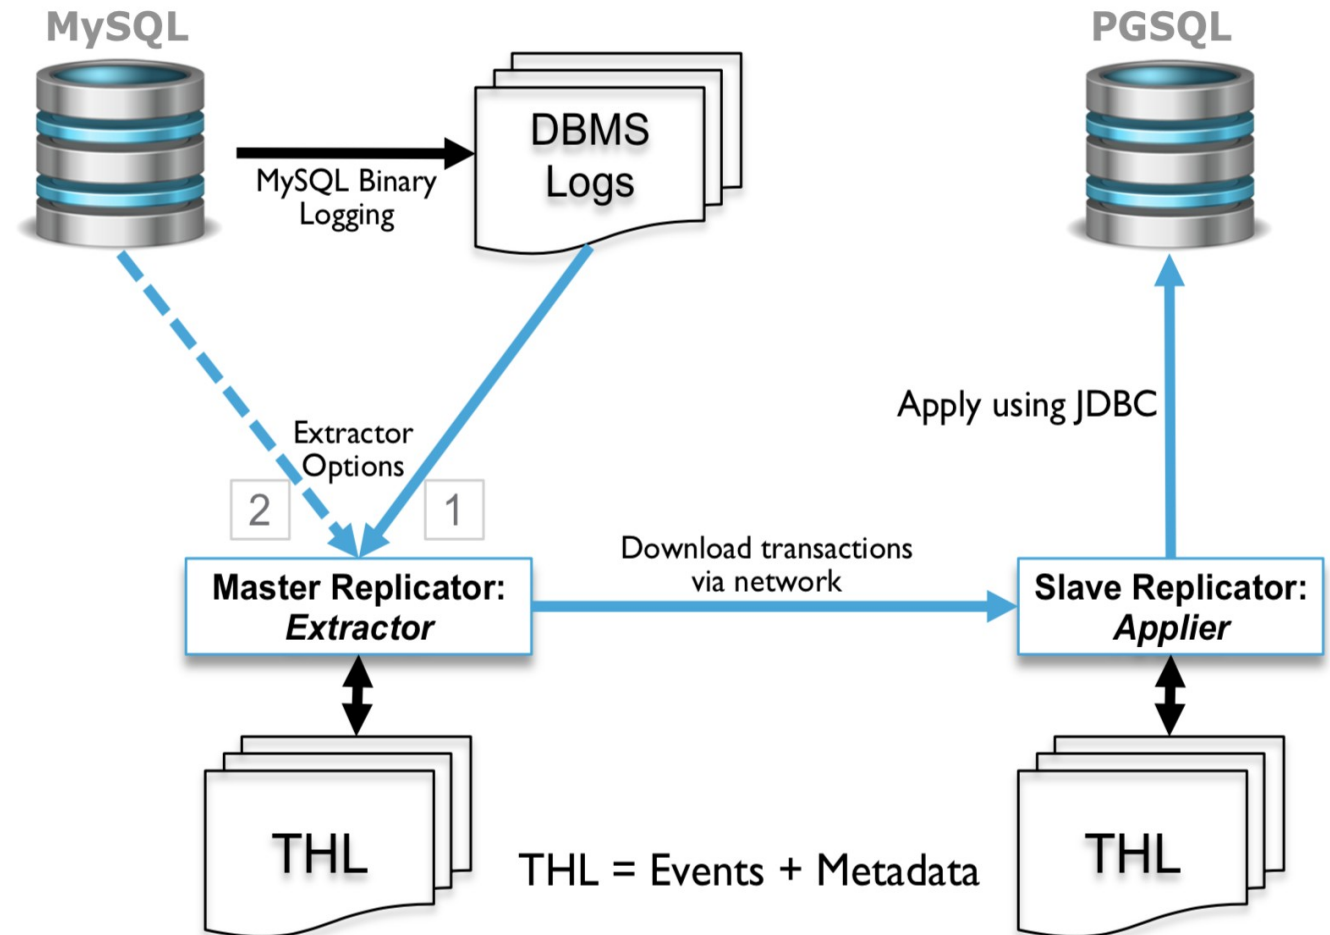
- Scalable;
- Asynchronous;
- Trigger based;
- One way and multi-master;
- Routing;
- Filtered Synchronization;
- Transformation Capabilities;



Available solutions...

Tungsten Replicator

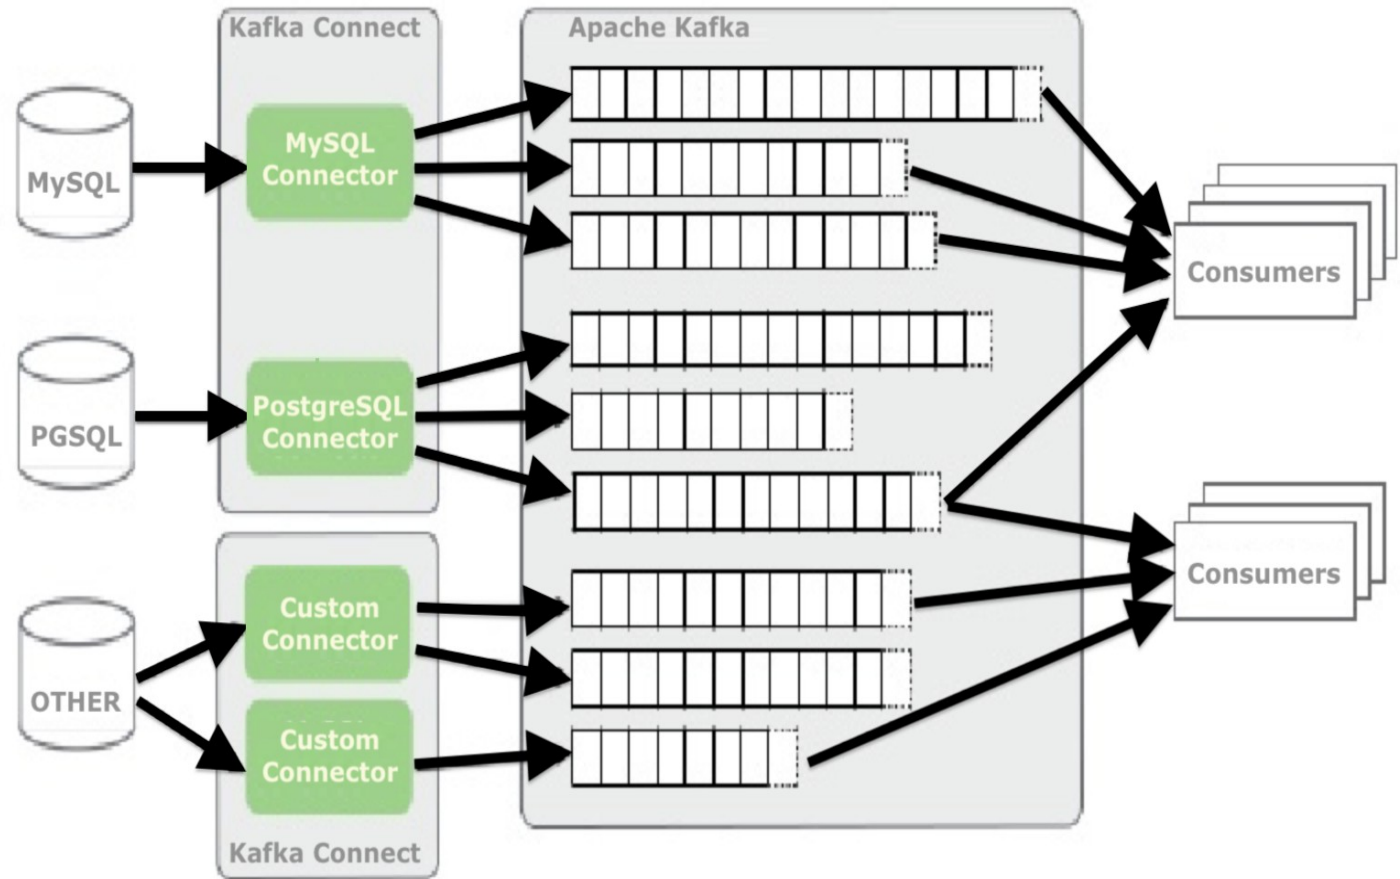
- Row-based (heterogeneous)
- Asynchronous;
- Transactionally Consistent;
- Flexible Extraction;
- Filtered Synchronization;
- Transformation Capabilities;



Available solutions...

Debezium (Apache Kafka)

- Row-based;
- Asynchronous/Queue (topics);
- Transactionally Consistent;
- Flexible Extraction;
- Filtered Synchronization;
- Transformation Capabilities;



Challenges

There are many challenges when replicating from different DB's:

- Data Type Differences;
- Provisioning;
- Database vs. Schema;
- SQL Dialect Differences;

Demo



Database Performance Matters